

Embedded Linux on Zynq UltraScale+ MPSoC Supplement

Tool : Xilinx Vivado & Petalinux 2019.2

Kit : Ultra96 Training Kit

What is Petalinux?

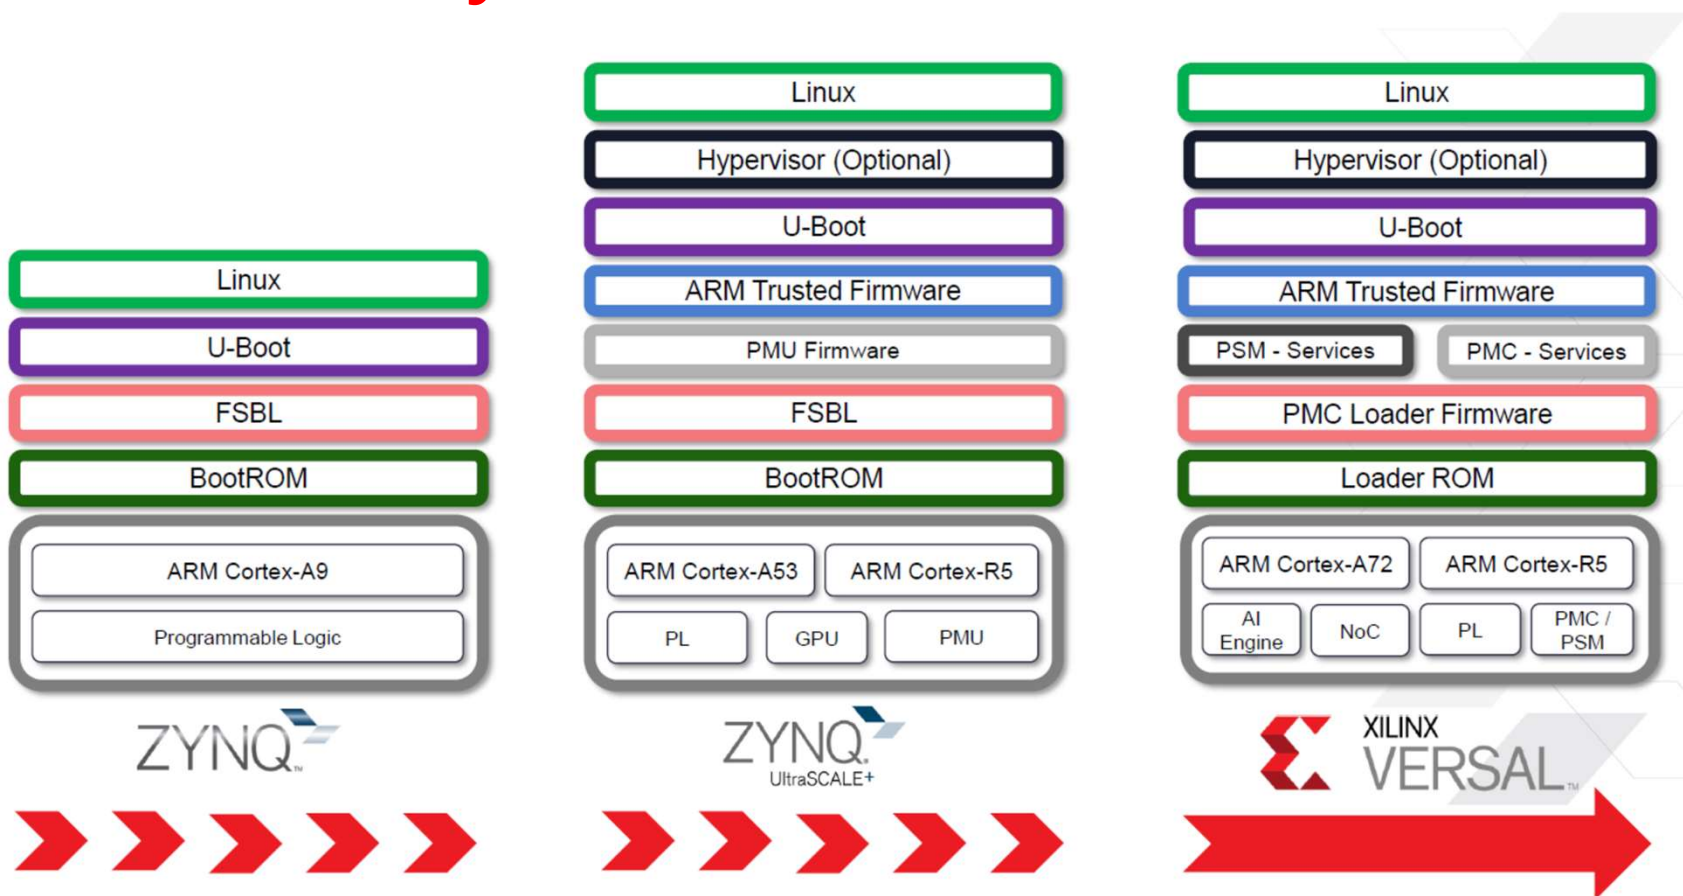
❖ Embedded Linux System Builder for Xilinx Boards

- ❖ A set of open-source cross-platform development command line and menu tools for use running under an x86 Linux Host OS
- ❖ Just six Petalinux commands (with many options):
petalinux-create, petalinux-config, petalinux-build, petalinux-package, petalinux-boot, petalinux-util
- ❖ Petalinux tools now use the open-source community Yocto system since v2016.4

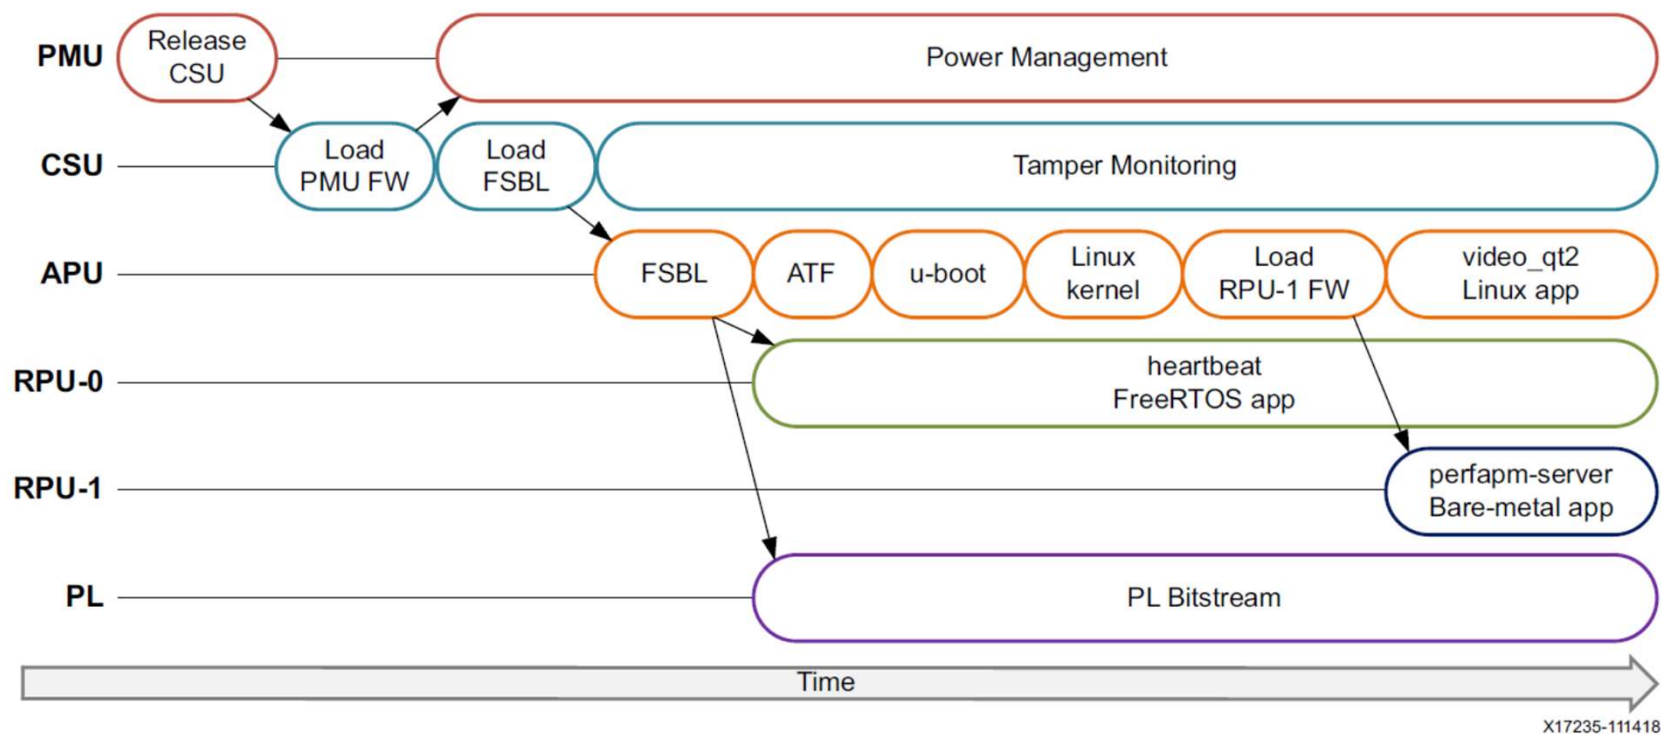
❖ Has Yocto recipes for BSPs, Libraries, Applications

- ❖ Mostly open-source based with a few exceptions for some hardware drivers
 - ❖ Petalinux includes and manages the Linux kernel sources and libraries
 - ❖ Various hardware drivers and modules
 - ❖ Conventional Linux applications and utilities
- ❖ Through Yocto, Petalinux can make use of additional embedded Linux applications beyond what Xilinx supplies

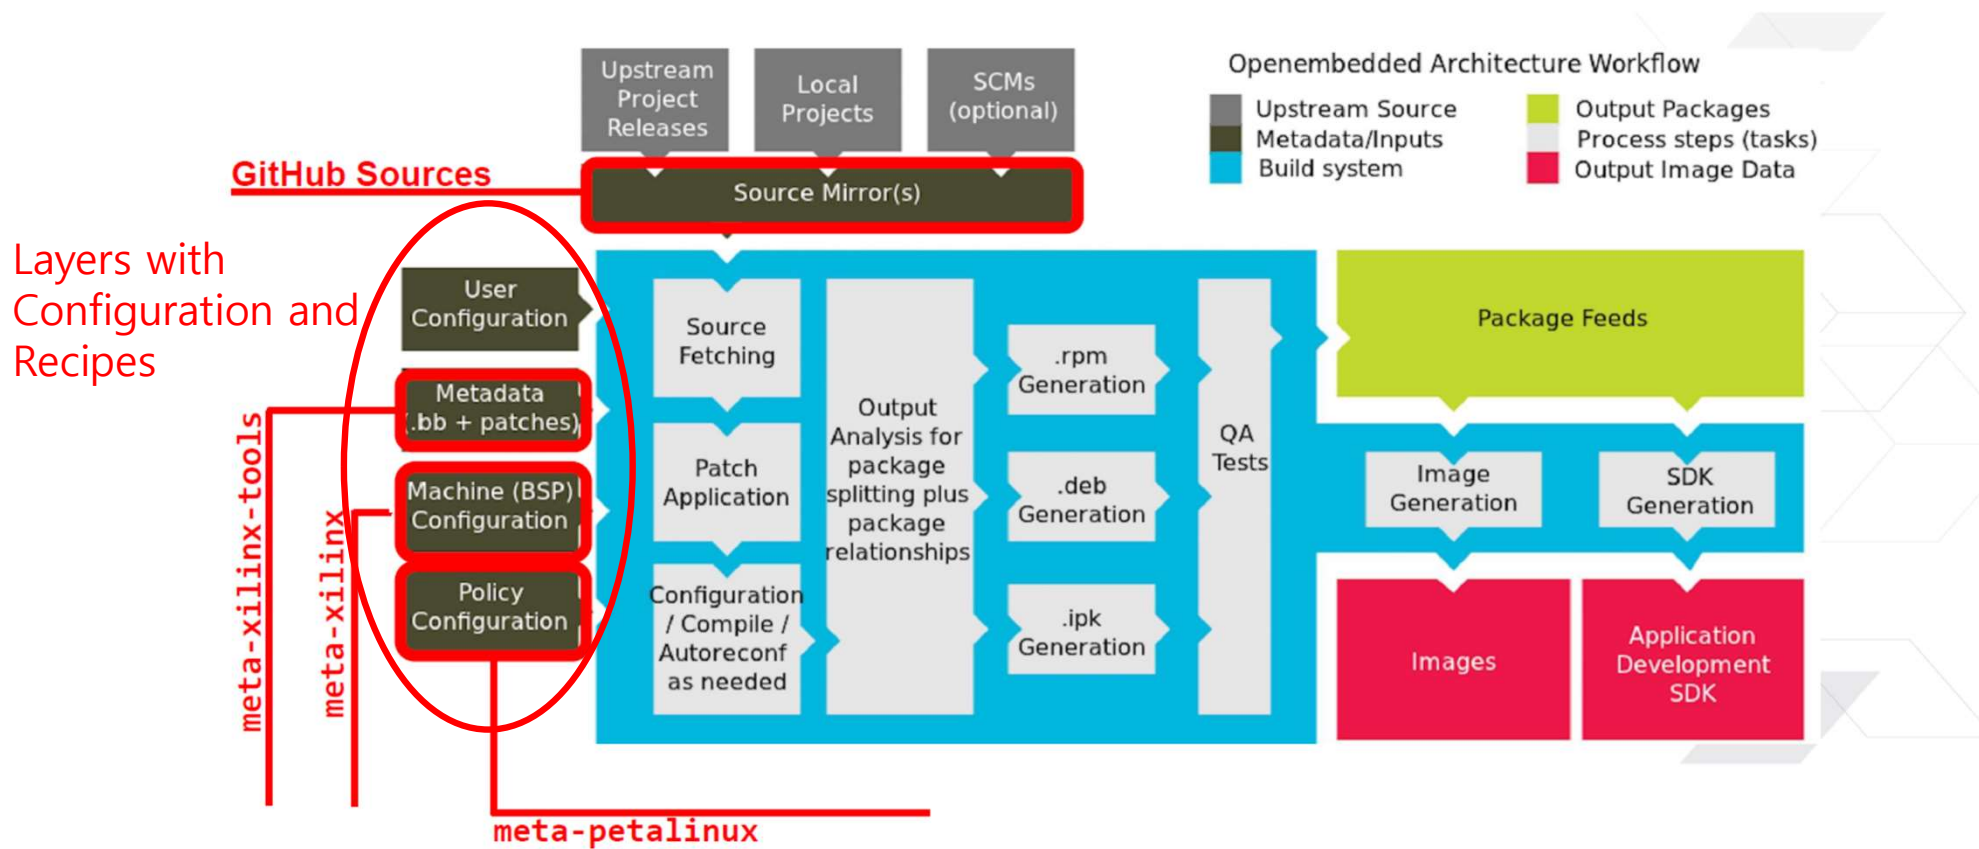
Products built by Petalinux



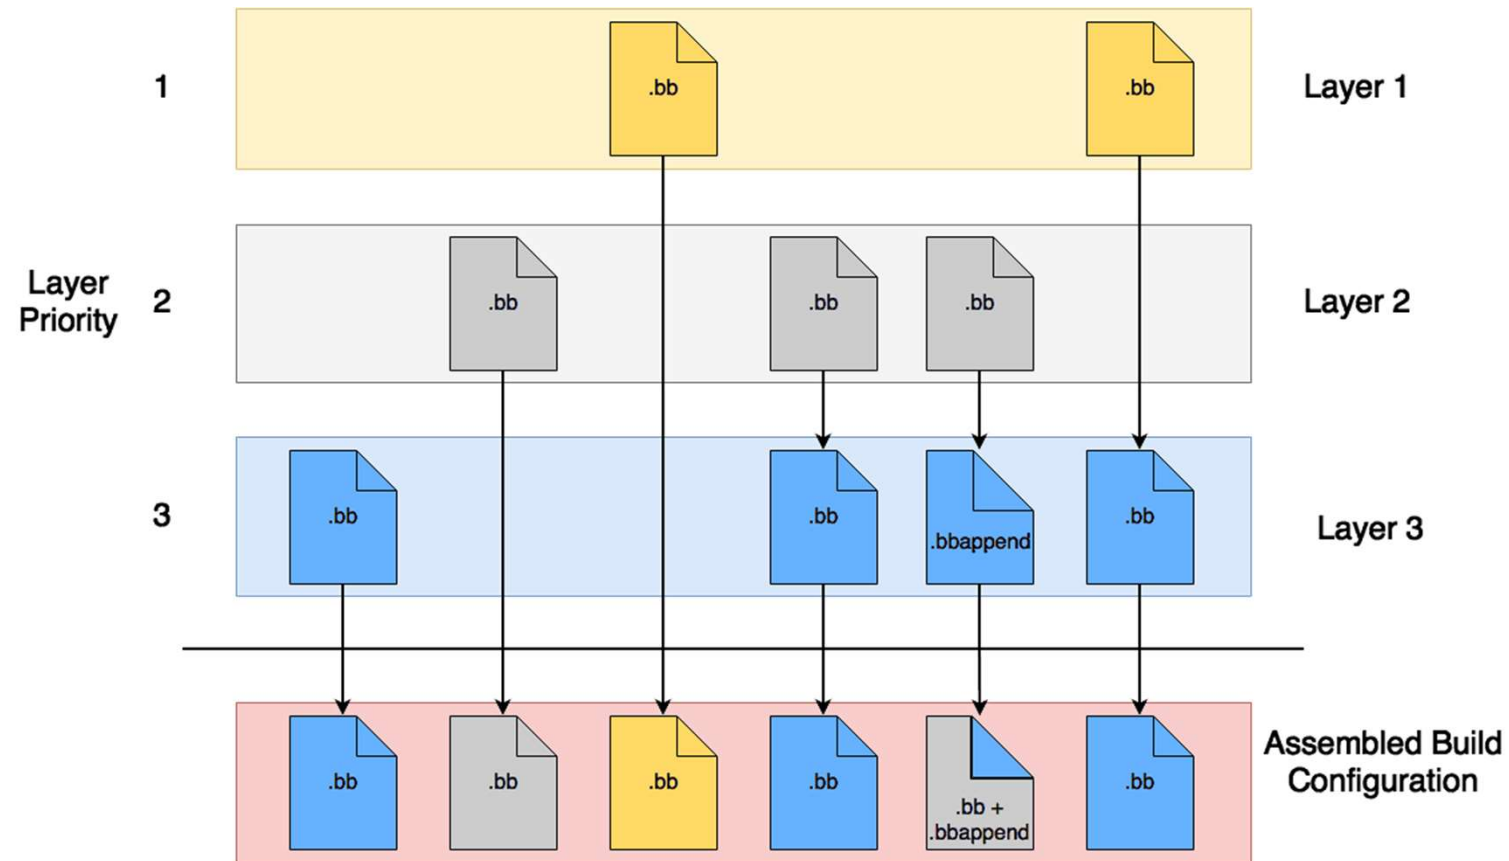
Products built by Petalinux



Petalinux Build Process



Layers with Recipes



Layers

- ❖ Petalinux build system is composed of layers
- ❖ Layers are containers for the recipes used to construct the system
- ❖ Layers are a way to manage extensions and customizations to the system
 - ❖ Layers can extend, add, replace or modify recipes
 - ❖ Layers are added via BBLAYERS variable in build/conf/bblayers.conf
 - ❖ petalinux-create generates meta-plnx-generated and meta-user layers
 - ❖ petalinux-config can add user layers(meta-inipro, meta-custom)

Layers

```
meta-inipro
├── conf
│   └── layer.conf
├── recipes-apps
│   ├── vai
│   │   ├── files
│   │   └── vai.bb
│   └── xrtutils
│       ├── files
│       └── xrtutils.bb
├── recipes-bsp
│   ├── u-boot
│   │   ├── u-boot-xlnx
│   │   └── u-boot-xlnx_2019.2.bbappend
│   ├── ultra96-misc
│   │   ├── files
│   │   └── ultra96-misc.bb
│   └── wilc-firmware
│       ├── wilc-firmware_15.00.bb
│       ├── wilc-firmware_15.01.bb
│       └── wilc-firmware_15.2.bb
└── recipes-connectivity
    ├── wpa-supPLICANT
    │   ├── files
    │   └── wpa-supPLICANT %.bbappend
```

bblayers.conf

```
LCONF_VERSION = "7"

BBPATH = "${TOPDIR}"
SDKBASEMETAPATH = "/media/hokim/data/petalinux/2019.2/components/yocto/source/aarch64"
BBLAYERS := " \
    ${SDKBASEMETAPATH}/layers/core/meta \
    ${SDKBASEMETAPATH}/layers/core/meta-poky \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-perl \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-python \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-filesystems \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-gnome \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-multimedia \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-networking \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-webserver \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-xfce \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-initramfs \
    ${SDKBASEMETAPATH}/layers/meta-openembedded/meta-oe \
    ${SDKBASEMETAPATH}/layers/meta-browser \
    ${SDKBASEMETAPATH}/layers/meta-qt5 \
    ${SDKBASEMETAPATH}/layers/meta-xilinx/meta-xilinx-bsp \
    ${SDKBASEMETAPATH}/layers/meta-xilinx/meta-xilinx-pynq \
    ${SDKBASEMETAPATH}/layers/meta-xilinx/meta-xilinx-contrib \
    ${SDKBASEMETAPATH}/layers/meta-xilinx-tools \
    ${SDKBASEMETAPATH}/layers/meta-petalinux \
    ${SDKBASEMETAPATH}/layers/meta-virtualization \
    ${SDKBASEMETAPATH}/layers/meta-openamp \
    ${SDKBASEMETAPATH}/layers/meta-jupyter \
    ${SDKBASEMETAPATH}/workspace \
    /home/hokim/work/zynqmp_linux/petalinux/ultra96/project-spec/meta-plnx-generated \
    /home/hokim/work/zynqmp_linux/petalinux/ultra96/project-spec/meta-user \
    /home/hokim/work/zynqmp_linux/petalinux/meta-inipro \
    /home/hokim/work/zynqmp_linux/petalinux/meta-custom \
    /home/hokim/work/zynqmp_linux/petalinux/ultra96/components/plnx_workspace \
"
```


Recipes

- ❖ Recipes for building packages

- ❖ Recipe Files

 - ❖ meta/recipes-devtools/dnf/dnf_2.7.5.bb

- ❖ Patches and Supplemental Files

 - ❖ Location : meta/recipes-devtools/dnf/dnf/

- ❖ Recipes inherit the system configuration and adjust it to describe how to build and adjust it to describe how to build and package the software

- ❖ Recipes can be extended and enhanced through append-files from other layers

Recipes

```
SUMMARY = "Package manager forked from Yum, using libsolv as a dependency resolver"
LICENSE = "GPLv2"
LIC_FILES_CHKSUM = "file://COPYING;md5=b234ee4d69f5fce4486a80fdaf4a4263 \
                    file://PACKAGE-LICENSING;md5=bfc29916e11321be06924c4fb096fdc"

SRC_URI = "git://github.com/rpm-software-management/dnf.git \
          file://0001-Correctly-install-tmpfiles.d-configuration.patch \
          file://0001-Do-not-hardcode-etc-and-systemd-unit-directories.patch \
          file://0005-Do-not-prepend-installroot-to-logdir.patch \
          file://0029-Do-not-set-PYTHON_INSTALL_DIR-by-running-python.patch \
          file://0030-Run-python-scripts-using-env.patch \
          "

SRCREV = "564c44667c7014843fa6f1732621093114ec59b2"
UPSTREAM_CHECK_GITTAGREGEX = "(?P<pver>\d+(\.\d+)+)"

S = "${WORKDIR}/git"

inherit cmake gettext bash-completion distutils3-base systemd

DEPENDS += "libdnf librepo libcomps python3-iniparse"

# manpages generation requires http://www.sphinx-doc.org/
EXTRA_OEMAKE = " -DWITH_MAN=0 -DPYTHON_INSTALL_DIR=${PYTHON_SITEPACKAGES_DIR} -DPYTHON_DESIRED=3"

BBCLASSEXTEND = "native nativesdk"

RDEPENDS_${PN}_class-target += " \
python3-core \
python3-codecs \
python3-netclient \
python3-email \
python3-threading \
python3-distutils \
python3-logging \
python3-fcntl \
librepo \
"
```

```
python3-shell \
libcomps \
libdnf \
python3-sqlite3 \
python3-compression \
python3-rpm \
python3-iniparse \
python3-json \
python3-curses \
python3-misc \
python3-gpg \
"

RRECOMMENDS_${PN}_class-target += "gnupg"

# Create a symlink called 'dnf' as 'make install' does not do it, but
# .spec file in dnf source tree does (and then Fedora and dnf documentation
# says that dnf binary is plain 'dnf').
do_install_append() {
    ln -s ${D}/${bindir}/dnf-3 ${D}/${bindir}/dnf
    ln -s ${D}/${bindir}/dnf-automatic-3 ${D}/${bindir}/dnf-automatic
}

# Direct dnf-native to read rpm configuration from our sysroot, not the one it was compiled in
do_install_append_class-native() {
    create_wrapper ${D}/${bindir}/dnf \
        RPM_CONFIGDIR=${STAGING_LIBDIR_NATIVE}/rpm \
        RPM_NO_CHROOT_FOR_SCRIPTS=1
}

SYSTEMD_SERVICE_${PN} = "dnf-makecache.service dnf-makecache.timer \
                        dnf-automatic.service dnf-automatic.timer \
                        dnf-automatic-download.service dnf-automatic-download.timer \
                        dnf-automatic-install.service dnf-automatic-install.timer \
                        dnf-automatic-notifyonly.service dnf-automatic-notifyonly.timer \
                        "

SYSTEMD_AUTO_ENABLE ?= "disable"
```

Device Tree

- ❖ The principle of Device Tree is to separate a large part of the hardware description from the kernel sources
- ❖ Device Tree allows a single kernel image to run on different boards with the differences being described in the device tree
- ❖ Device Tree is a tree of nodes that models the hierarchy of devices in the system, from the devices inside the processor to the devices on the board
- ❖ Each node can have a number of properties describing various properties of the devices: addresses, interrupts, clocks, etc.
- ❖ Written in a specialized language, the Device Tree source code is compiled into Device Tree Blob by the Device Tree Compiler(DTC)

Device Tree

- ❖ The DTC checks the device tree syntax but the semantics of the device tree are checked at runtime by the kernel and drivers
- ❖ At boot time, the kernel is given a compiled device tree, referred to as a Device Tree Blob, which is parsed to instantiate all the devices described in the device tree
- ❖ Some key properties in a device tree node
 - ❖ The **compatible** property is used to bind a device with a device driver
 - ❖ The **interrupts** property contains the interrupt number used by the device
 - ❖ The **reg** property contains the memory range used by the device
- ❖ There is limited documentation for the device tree bindings for each device such that driver code inspection may be necessary
 - ❖ The docs are in the kernel tree at Documentation/devicetree/bindings

Device Tree – A Simple Example

- ❖ A simple example below illustrates a node of a device
 - ❖ SomeAn AMBA bus with a GPIO that has registers mapped to 0x4120000 and is using interrupt 91
91 – 32 = 59, where 32 is the first Shared Peripheral Interrupt
 - ❖ The device is compatible with a driver containing a matching compatible string of **"xlnx, simple"**
 - ❖ The device driver source code may be the only way to really understand what properties it is expecting from the device tree

```
ps7_axi_interconnect_0: amba@0 {  
    #address-cells = <1>;  
    #size-cells = <1>;  
    compatible = "xlnx,ps7-axi-interconnect-1.00.a", "simple-bus";  
    ranges ;  
    axi_gpio_0: gpio@41200000 {  
        #gpio-cells = <2>;  
        compatible = "xlnx, simple";  
        gpio-controller ;  
        interrupt-parent = <&ps7_scugic_0>;  
        interrupts = <0 59 4>;  
        reg = <0x41200000 0x10000>;  
        xlnx,is-dual = <0x1>;  
    } ;  
};
```

Device Tree – Inclusion Example

ps.dtsi (included file)

```
ps7_ttc_1: ps7-ttc@0xf8002000 {  
    clocks = <&clkc 6>;  
    compatible = "xlnx,ps7-ttc-1.00.a";  
    interrupt-parent = <&ps7_scugic_0>;  
    interrupts = <0 37 4>, <0 38 4>, <0 39 4>;  
    reg = <0xF8002000 0x1000>;  
    status = "disabled";  
};
```

system-top.dts (including file)

```
/include/ "ps.dtsi"  
  
&ps7_ttc_1 {  
    compatible = "xlnx,psttc", "generic-uio";  
    status = "okay";  
};
```

+

=

```
ps7_ttc_1: ps7-ttc@0xf8002000 {  
    clocks = <&clkc 6>;  
    compatible = "xlnx,psttc", "generic-uio";  
    interrupt-parent = <&ps7_scugic_0>;  
    interrupts = <0 37 4>, <0 38 4>, <0 39 4>;  
    reg = <0xF8002000 0x1000>;  
    status = "okay";  
};
```

Note the "&" used to reference an existing node (rather than creating a new node)

The result for the duplicated (red) properties is the same as the including file.

Device Tree – Simplified Ultra96v2

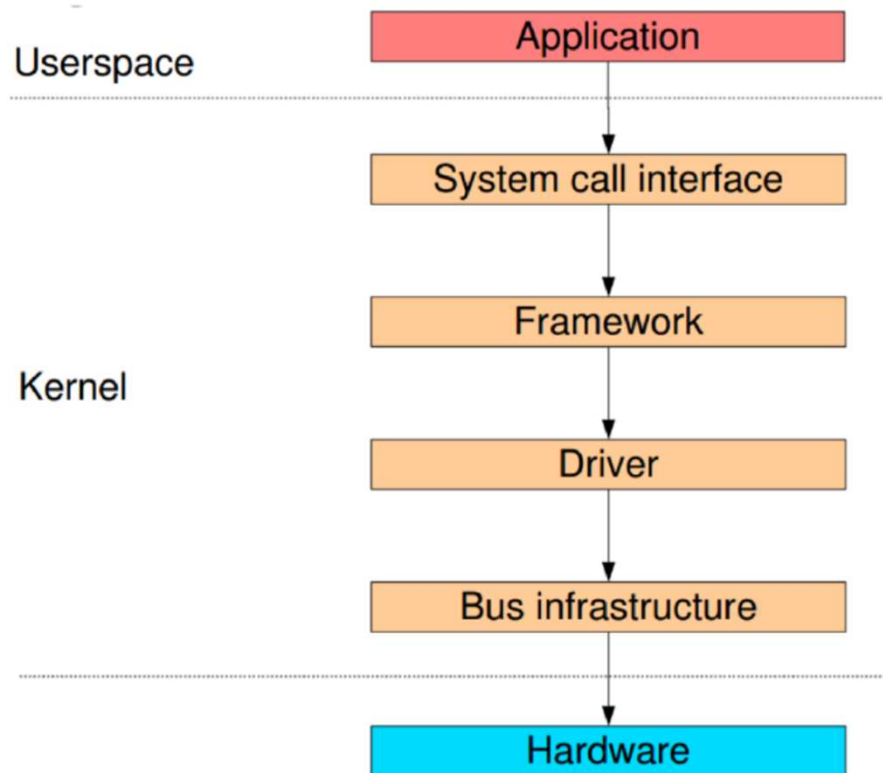
```
/ {
    compatible = "avnet,ultra96-rev1", "avnet,ultra96", "xlnx,zynqmp-zcu100-revC",
        "xlnx,zynqmp-zcu100", "xlnx,zynqmp";
    #address-cells = <0x2>;
    #size-cells = <0x2>;
    model = "Avnet Ultra96 Rev1";
    cpus {
        cpu@0 {
            compatible = "arm,cortex-a53", "arm,armv8";
        };
    };
    amba {
        compatible = "simple-bus";
        serial@ff010000 {
            compatible = "cdns,uart-rlp12", "xlnx,xuartps";
        };
        i2c@ff030000 {
            compatible = "cdns,i2c-rlp14", "cdns,i2c-rlp10";
            i2c-mux@75 {
                compatible = "nxp,pca9548";
                i2c@4 {
                    reg = <0x4>;
                    irps5401@13 {
                        compatible = "infineon,irps5401";
                        reg = <0x13>;
                    };
                };
            };
        };
    };
    mmc@ff170000 {
        compatible = "xlnx,zynqmp-8.9a", "arasan,sdhci-8.9a";
        wilc_sdio@1 {
            compatible = "microchip,wilc3000";
        };
    };
};
```

```
amba_pl@0 {
    compatible = "simple-bus";
    gpio@80000000 {
        compatible = "xlnx,axi-gpio-2.0", "xlnx,xps-gpio-1.00.a";
    };
    i2c@80010000 {
        compatible = "xlnx,axi-iic-2.0", "xlnx,xps-iic-2.00.a";
        pmodtmp2@4b {
            compatible = "inipro,pmodtmp2";
        };
    };
    axi_quad_spi@80020000 {
        compatible = "xlnx,axi-quad-spi-3.2", "xlnx,xps-spi-2.00.a";
        pmodals@0 {
            compatible = "inipro,pmodals";
        };
    };
};
chosen {
    bootargs = "earlycon console=ttyPS0,115200 clk_ignore_unused root=/dev/m
mcblk0p2 rw rootwait uio_pdrv_genirq.of_id=xlnx,generic-uio cma=512M";
    stdout-path = "serial0:115200n8";
};
aliases {
    serial0 = "/amba/serial@ff010000";
};
memory {
    device_type = "memory";
    reg = <0x0 0x0 0x0 0x7ff00000>;
};
};
```


Linux Kernel Frameworks

- ❖ Many device drivers are not directly implemented as character devices or block devices. They are implemented under a framework, specific to a device type (framebuffer, V4L, serial, etc.)
 - ❖ The framework factors out the common parts of drivers for the same type of devices to reduce code duplication
 - ❖ From userspace, many are still seen as normal character devices
 - ❖ The frameworks provide a coherent userspace interface (ioctl numbering and semantics, etc.) for every type of device, regardless of the driver
 - ❖ The network framework of Linux provides a socket API such that an application can connect to a network using any network driver without knowing the details of the network driver
- ```
sockfd = socket(AF_INET, SOCK_STREAM, 0);
```

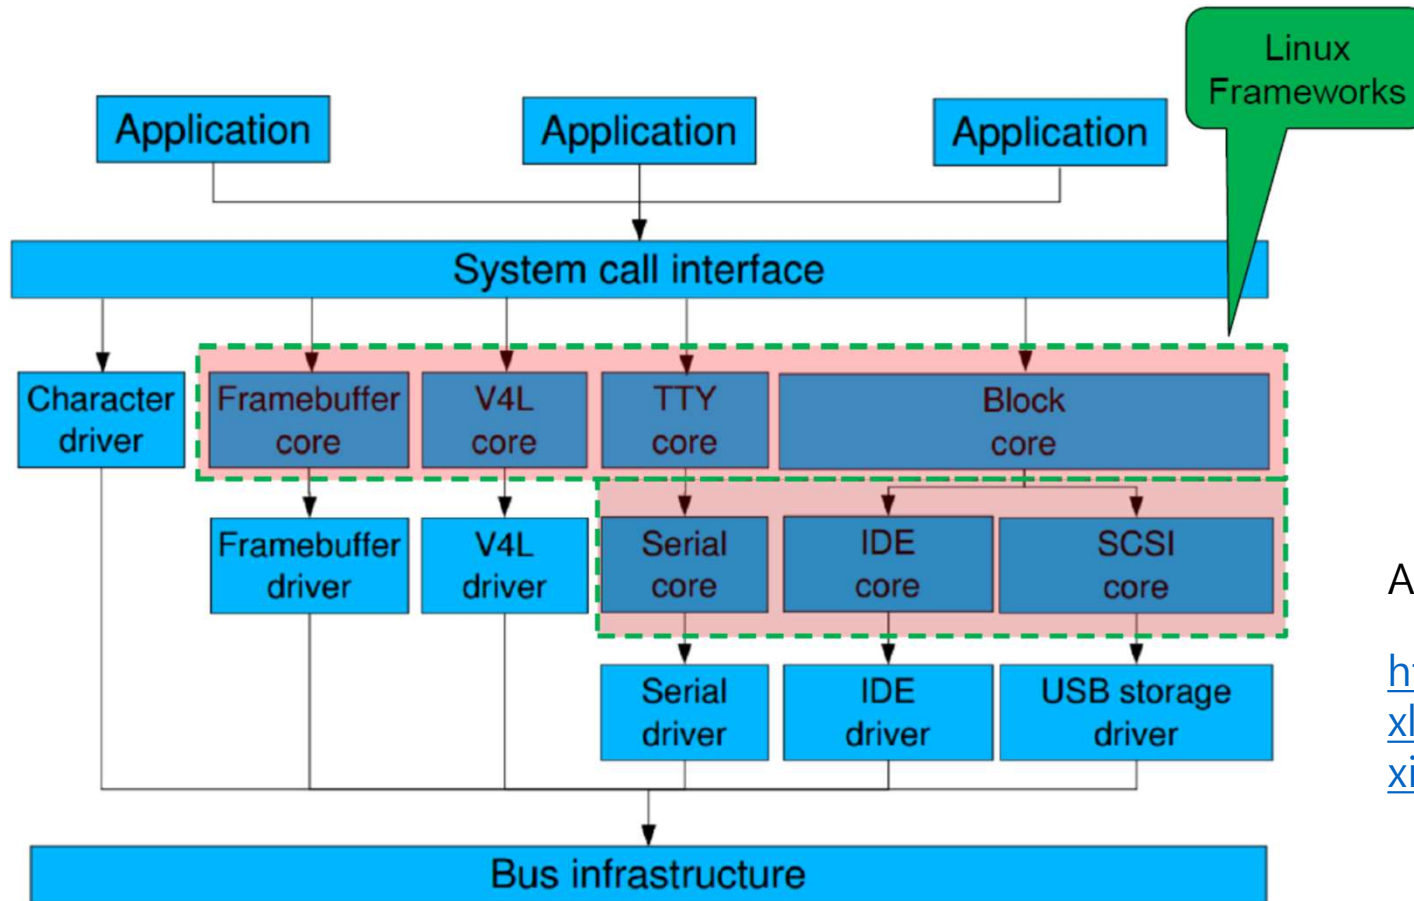
# Linux Kernel Layers Focused on Frameworks



❖ A driver is always interfacing with:

- ❖ A framework that allows the driver to expose the hardware features to userspace applications
- ❖ A bus infrastructure(part of the device model), to detect/communicate with the hardware

# Example Frameworks



An example driver of spi framework

<https://github.com/Xilinx/linux-xlnx/blob/master/drivers/spi/spi-xilinx.c#L846>

# Virtual File Systems - Overview

- ❖ System and kernel information

- ❖ Presented to userspace application as virtual file systems
- ❖ Created dynamically and only exist in memory

- ❖ Two virtual filesystems most known to users

- ❖ proc, mounted on /proc, contains operating system related information (processes, memory management parameters...)

This is an older mechanism that became somewhat chaotic

- ❖ sysfs, mounted on /sys, contains a representation of the system as a set of devices and buses together with information about these devices

This is the newer mechanism and is the preferred place to add system information

# Virtual File Systems - sysfs

- ❖ The **sysfs** virtual filesystem is a mechanism for the kernel to export operating details to userspace
- ❖ The kernel exports the following items to userspace
  - ❖ The bus, device, drivers, etc. structures internal to the kernel
  - ❖ **/sys/bus/** contains the list of buses
  - ❖ **/sys/devices/** contains the list of devices
  - ❖ **/sys/class/** enumerates devices by class (net, input, block...), whatever the bus they are connected to

# Kernel Modules

- ❖ The Linux kernel by design is a monolithic kernel, but is also modular
- ❖ The kernel can dynamically load and unload parts of the kernel code which are referred to kernel modules
- ❖ Modules allow the kernel capabilities to be extended without modifying the rest of the code or rebooting the kernel
- ❖ A kernel module can be inserted or removed while the kernel is running
  - ❖ It can be inserted manually by a root user or from a user space script at startup
- ❖ Kernel modules help to keep the kernel size to a minimum and makes the kernel very flexible
- ❖ Kernel modules are useful to reduce boot time since time is not spent initializing devices and kernel features that are only needed later
- ❖ Once loaded, kernel modules have full control and privileges in the system such that only the root user can load and unload modules

# Kernel Modules Details

- ❖ Naming Conversion: <filename>.ko
- ❖ Location: /lib/modules/<kernel\_version> on the root filesystem
- ❖ Device drivers can be kernel modules or statically built into the kernel image
- ❖ Building a device driver as a module makes the development easier since it can be loaded, tested, and unloaded without rebooting the kernel

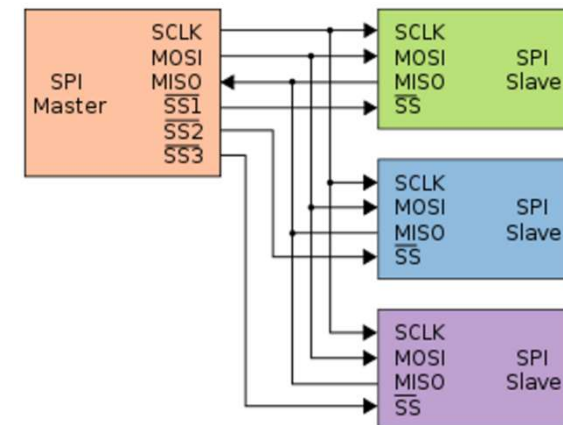
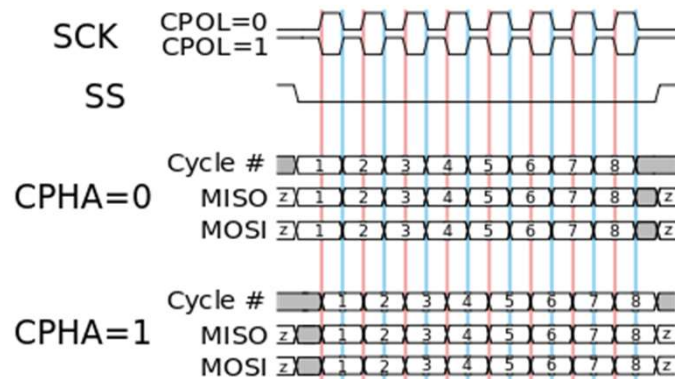


# SPI

❖ Synchronous serial digital data link by Motorola

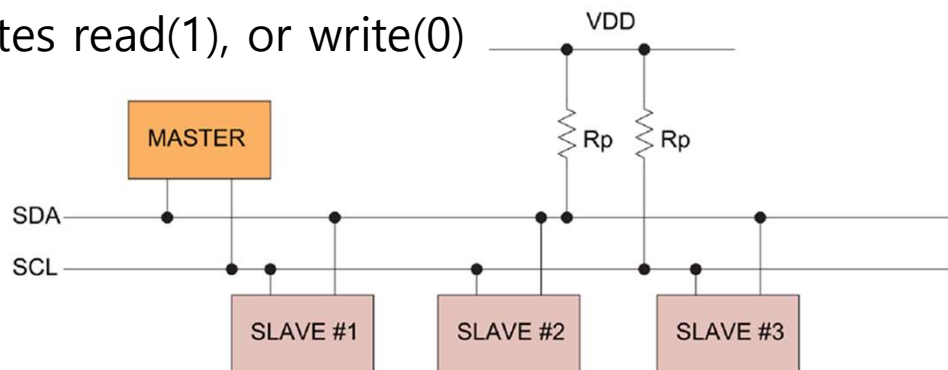
❖ SPI master controls

- ❖ CLK: synchronization clock
- ❖ SS or CS: slave select or chip-select; when active the slave is allowed to talk
- ❖ MOSI: master out, slave in; carries data from master to slave
- ❖ MISO: master in, slave out; carries data from slave to master



# I2C

- ❖ Two wires are controlled by the master and slaves according to the protocol
  - ❖ SCL: Serial Clock
  - ❖ SDA: Serial Data
- ❖ Each have has a 7 bit address
- ❖ The 7 MSB of the first transmitted byte are slave address
- ❖ The LSB of this byte indicates read(1), or write(0)



## Kernel UIO API – Sys Filesystem

- ❖ The UIO driver in the kernel creates file attributes in the sys filesystem describing the UIO device
- ❖ `/sys/class/uio` is the root directory for all the file attributes
- ❖ A separate numbered directory structure is created under `/sys/class/uio` for each UIO device
  - ❖ First UIO device: `/sys/class/uio/uio0`
  - ❖ `/sys/class/uio/uio0/name` contains the name of the device which correlates to the name in the `uio_info` structure
  - ❖ `/sys/class/uio/uio0/maps` is a directory that has all the memory ranges for the device
  - ❖ Each numbered map directory has attributes to describe the device memory including the address, name, offset and size (`/sys/class/uio/uio0/maps/map0`)

# User Space Driver Example

```
#define UIO_SIZE "/sys/class/uio/uio0/maps/map0/size"
```

```
int main(int argc, char **argv) {
```

```
 int uio_fd;
```

```
 unsigned int uio_size;
```

```
 FILE *size_fp;
```

```
 void *base_address;
```

```
 uio_fd = open("/dev/uio0", O_RDWR);
```

```
 size_fp = fopen(UIO_SIZE, O_RDONLY);
```

```
 fscanf(size_fp, "0x%08X", &uio_size);
```

```
 base_address = mmap(NULL, uio_size,
 PROT_READ | PROT_WRITE,
 MAP_SHARED, uio_fd, 0);
```

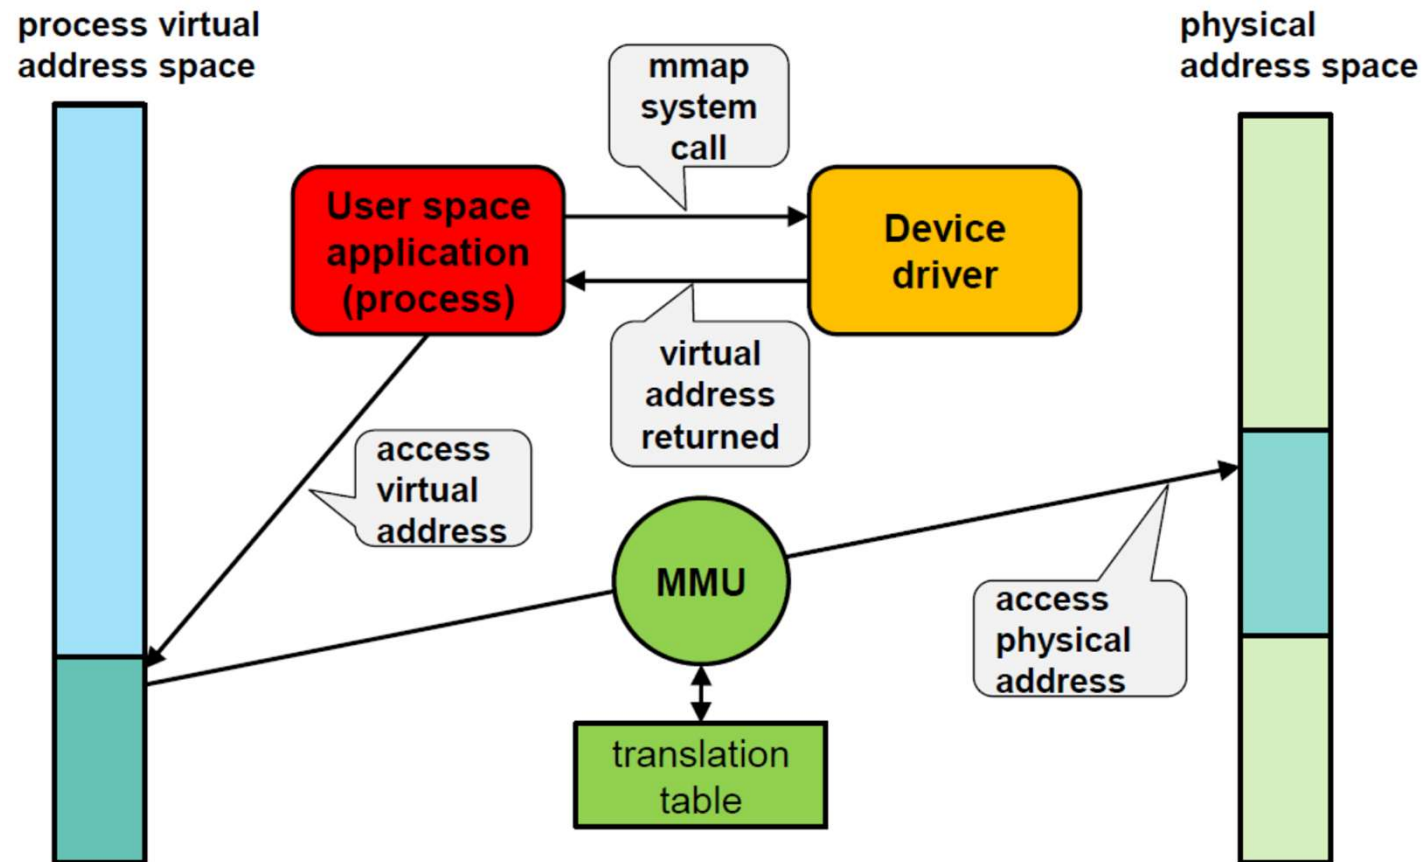
```
 // Access to the hardware can now occur....
```

```
 munmap(base_address, uio_size);
```

```
}
```

- ❖ Open the UIO device so that it's ready to use
- ❖ Get the size of the memory region from the size sysfs file attribute
- ❖ Map the device registers into the process address space so they are directly accessible
- ❖ Unmap the device registers to finish

# Mapping Device Memory Flow




# ZOCL Driver Interfaces

- ❖ A GEM style driver for Xilinx edge based accelerators
- ❖ Accelerator memory allocation is modeled as buffer objects(bo), supports both SMMU based shared virtual memory and **CMA based shared physical memory between PS and PL**. zocl also supports memory management of PL-DDRs and PL-BRAMs. PL-DDR is reserved by zocl driver via device tree. **Both PS Linux and PL logic can access PL-DDRs**
- ❖ ioctl command codes and associated structures for zocl driver

[https://xilinx.github.io/XRT/master/html/zocl\\_ioctl.main.html](https://xilinx.github.io/XRT/master/html/zocl_ioctl.main.html)

# Media pipeline for camera



```

ov5640: camera@3c {
 port {
 ov5640_out: endpoint {
 remote-endpoint = <&csiss_in>;
 };
 };
};

```

```

mipi_csi2_rx_subsystem_0: mipi_csi2_rx_subsystem@80000000 {
 csiss_ports: ports {
 csiss_port0: port@0 {
 csiss_out: endpoint {
 remote-endpoint = <&vcap_mipi_in>;
 };
 };
 csiss_port1: port@1 {
 csiss_in: endpoint {
 remote-endpoint = <&ov5640_out>;
 };
 };
 };
};

```

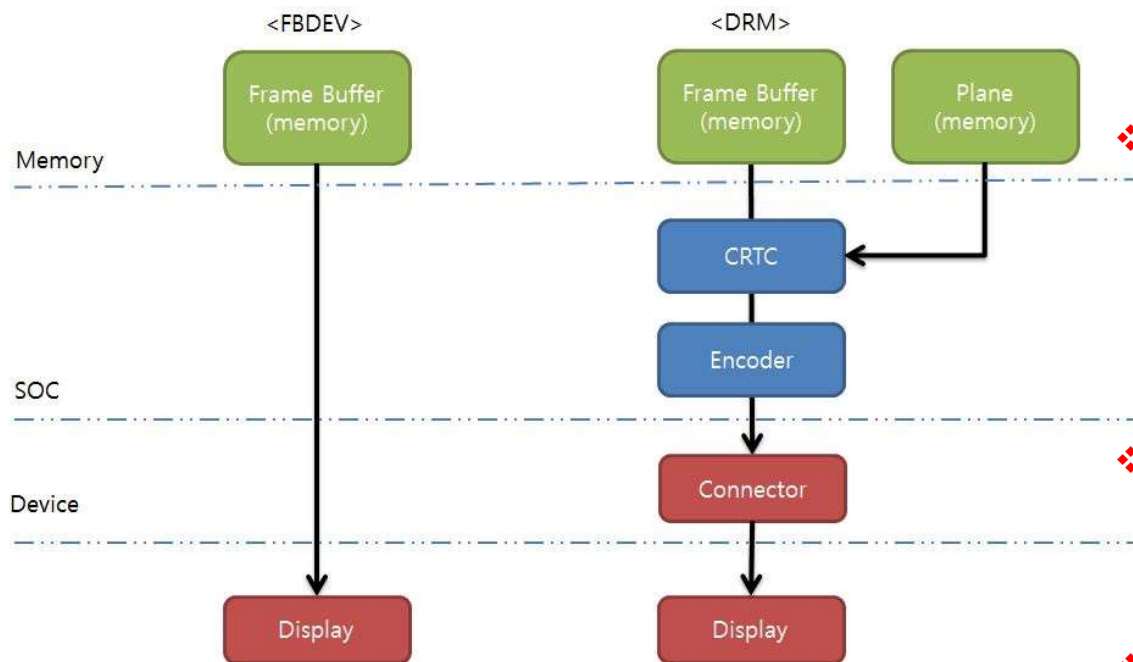
```

vcap_mipi{
 vcap_mipi_ports: ports {
 vcap_mipi_port: port@0 {
 vcap_mipi_in: endpoint {
 remote-endpoint = <&csiss_out>;
 };
 };
 };
};

```



# DRM Components



## ❖ Plane

- ❖ An image layer
- ❖ The final image displayed by the CRTC is the composition of one or several planes

## ❖ CRTC

- ❖ Mode information: resolution, depth, polarity, porch, refresh rate, and so on
- ❖ Information of the buffer region displayed
- ❖ Change current framebuffer to new one

## ❖ Encoder

- ❖ Take the digital bit-stream from the CRTC
- ❖ Convert to the appropriate analog levels

## ❖ Connector

- ❖ Provide the appropriate physical plugs such as HDMI, VGA, DP and so on